


Graph Neural Networks



From Images to Graphs



Aggregate feature information
from neighbors

Motivation: Extend Convolutional Neural Networks (CNNs) to Graphs 🤔





Image Processing Fundamentals: Edge Detection



Original Image



Edge Detected Image



Basic Image Processing



- Image = 2D matrix of pixel values
- Each pixel represents brightness/color
- Example grayscale image:

$$X = \begin{bmatrix} 10 & 10 & 80 & 10 & 10 & 10 \\ 10 & 10 & 80 & 10 & 10 & 10 \\ 10 & 10 & 80 & 10 & 10 & 10 \\ 10 & 10 & 80 & 10 & 10 & 10 \\ 10 & 10 & 80 & 10 & 10 & 10 \\ 10 & 10 & 80 & 10 & 10 & 10 \end{bmatrix}$$

57	153	174	168	150	152	129	151	172	161	155	166
56	182	163	74	35	62	33	17	110	210	180	164
80	180	50	14	94	6	10	33	46	105	159	181
206	109	5	124	191	111	120	204	166	15	55	180
94	65	157	251	237	259	239	238	227	87	71	201
72	106	267	233	233	214	230	230	228	95	74	206
88	83	179	209	185	215	211	158	139	75	33	161
89	97	165	84	10	168	134	11	31	62	22	145
99	168	191	163	158	227	178	143	182	105	36	190
106	174	155	252	236	231	149	178	228	43	95	234
90	216	116	149	236	187	85	150	79	35	218	241
90	234	147	108	227	210	127	102	36	101	255	224
90	214	179	65	109	143	95	50	2	109	249	219
87	196	235	75	1	81	47	0	6	217	255	211
80	202	237	145	0	0	12	108	200	138	243	236
95	206	123	207	177	121	120	200	175	13	96	218

Convolution: Spatial Domain

- Slide kernel over image
- Multiply and sum values
- Example kernel (vertical edge detection):
- Demo

$$K = \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$

7	2	3	3	8
4	5	3	8	4
3	3	2	8	4
2	8	7	2	7
5	4	4	5	4

$$\begin{array}{c} * \\ \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix} \\ = \\ \begin{bmatrix} 6 & & \\ & & \\ & & \end{bmatrix} \end{array}$$

$$\begin{aligned} 7x1 + 4x1 + 3x1 + \\ 2x0 + 5x0 + 3x0 + \\ 3x-1 + 3x-1 + 2x-1 \\ = 6 \end{aligned}$$

Convolution is Complicated 😬

Example:

Suppose we have an image X and a kernel K as follows:

$$X = [X_1 \quad X_2 \quad X_3 \quad X_4 \quad X_5 \quad X_6]$$
$$K = [K_1 \quad K_2 \quad K_3]$$

The convolution is given by

$$X * K = \sum_{i=1}^6 X_i K_{7-i}$$

Or equivalently,

$$X * K = [X_1 K_3 + X_2 K_2 + X_3 K_1 \quad X_2 K_3 + X_3 K_2 + X_4 K_1 \quad X_3 K_3 + X_4 K_2 + X_5 K_1 \quad X_4 K_3 + X_5 K_2 + X_6 K_1]$$

Let's make it simpler using the convolution theorem!

What is the convolution theorem?

Suppose two functions f and g and their Fourier transforms F and G . Then,

$$\underbrace{(f * g)}_{\text{convolution}} \leftrightarrow \underbrace{(F \cdot G)}_{\text{multiplication}}$$

The Fourier transform is a one-to-one mapping between f and F (and g and G).

But what is the Fourier transform 😊?


Fourier Transform: The Basics

Transform a signal from:

- Time/Space domain →
Frequency domain

Key Concept:

- Any signal can be decomposed
into sum of sine and cosine waves
- Each wave has specific frequency
and amplitude




2D Fourier Transform



2D Fourier Transform decomposes image into sum of $2D$ waves.

$$\mathcal{F}(X)[h, w] = \sum_{k=0}^{H-1} \sum_{\ell=0}^{W-1} X[k, \ell] \cdot \underbrace{e^{-2\pi i \left(\frac{hk}{H} + \frac{w\ell}{W} \right)}}_{2D \text{ wave}}$$


For image X with size $H \times W$.



Edge Detection in Frequency Domain



Original Image → Fourier Transform → Apply Filter → Inverse Transform




JPEG Compression



1. Divide image into 8x8 blocks
2. Apply Discrete Cosine Transform (similar to Fourier)
3. Quantize frequencies
 - Keep low frequencies
 - Discard high frequencies
4. Encode efficiently

Benefits:

- Smaller file size
- Maintains visual quality
- Exploits human visual perception





Coding Exercise

Coding Exercise

From Images to Graphs

- Image = 2D grid of pixels
- Through a convolution, a pixel value is influenced by its neighbors
- We can represent this neighborhood structure using a graph and define **convolutions on graphs!**



Graph Fourier Transform



- Just like signals, graphs have a "frequency" domain.
- Suppose we have a graph of N nodes, each node has a feature x_i .
- **The total variation** measures the smoothness of the node features:

$$J = \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N A_{ij}(x_i - x_j)^2 = \mathbf{x}^\top \mathbf{L} \mathbf{x}$$

where \mathbf{L} : Graph Laplacian, x_i : Node features, A_{ij} : Adjacency matrix

Q: What x makes the total variation smallest (most smooth) and largest (most varying)? 🤔

The eigendecomposition of the Laplacian:

$$\mathbf{L}\mathbf{x} = \lambda\mathbf{x}$$

By multiplying both sides by \mathbf{x}^\top , we get

$$\mathbf{x}^\top \mathbf{L} \mathbf{x} = \lambda$$

This tells us that:

1. The eigenvectors with small eigenvalues represent **low-frequency** signals.
2. The eigenvectors with large eigenvalues represent **high-frequency** signals.

Decomposing the Total Variation

The total variation can be decomposed as follows (\mathbf{u}_i is the eigenvector of the Laplacian):

$$\begin{aligned} J &= \mathbf{x}^\top \mathbf{L} \mathbf{x} = \mathbf{x}^\top \left(\sum_{i=1}^N \lambda_i \mathbf{u}_i \mathbf{u}_i^\top \right) \mathbf{x} = \sum_{i=1}^N \lambda_i (\mathbf{x}^\top \mathbf{u}_i)(\mathbf{u}_i^\top \mathbf{x}) \\ &= \sum_{i=1}^N \lambda_i \underbrace{\|\mathbf{x}^\top \mathbf{u}_i\|^2}_{\text{alignment between } \mathbf{x} \text{ and } \mathbf{u}_i} \end{aligned}$$

Key Insight:

- The total variation is now decomposed into the sum of different frequency components $\lambda_i \cdot \|\mathbf{x}^\top \mathbf{u}_i\|^2$.
- λ_i acts as a *filter (kernel)* that reinforces or passes the signal $\mathbf{x}^\top \mathbf{u}_i$.

Spectral Filtering



Low-pass Filter:

$$h_{\text{low}}(\lambda) = \frac{1}{1 + \alpha\lambda}$$

High-pass Filter:

$$h_{\text{high}}(\lambda) = \frac{\alpha\lambda}{1 + \alpha\lambda}$$

Properties:

- Controls which frequencies pass
- Smooth or sharpen signals

Spectral Graph Convolution



Key idea: Let the kernel be a learnable parameter:

$$\mathbf{L}_{\text{learn}} = \sum_{k=1}^K \theta_k \mathbf{u}_k \mathbf{u}_k^\top$$

Layer operation:

$$\mathbf{x}^{(\ell+1)} = h \left(\mathbf{L}_{\text{learn}} \mathbf{x}^{(\ell)} \right)$$

Where:

- θ_k : Learnable parameters
- h : Activation function
- K : Number of filters
- <https://arxiv.org/abs/1312.6203>

Multi-dimensional Features

Let's consider the case where the node features are multi-dimensional. We want to map the f_{in} -dimensional features to f_{out} -dimensional features.

Key idea: Learn a separate filter for every combination of individual input and output features.

$$\mathbf{X}'_i = h \left(\sum_{j=1}^{f_{\text{in}}} L_{\text{learn}}^{(i,j)} \mathbf{X}_j^{(\ell)} \right), \quad \forall i \in \{1, \dots, f_{\text{out}}\}$$

Where:

- $\mathbf{X} \in \mathbb{R}^{N \times f_{\text{in}}}$
- $\mathbf{X}' \in \mathbb{R}^{N \times f_{\text{out}}}$
- $L_{\text{learn}}^{(i,j)} = \sum_{k=1}^K \theta_{k,(i,j)} \mathbf{u}_k \mathbf{u}_k^\top$

Limitations of Spectral GNNs !

1. Computational Cost:

- Eigendecomposition: $O(N^3)$
- Prohibitive for large graphs

2. Spatial Locality:

- Non-localized filters
- Global node influence

ChebNet: A Bridge Solution



Key Idea: Approximate filters using Chebyshev polynomials

$$\mathbf{L}_{\text{learn}} \approx \sum_{k=0}^{K-1} \theta_k T_k(\tilde{\mathbf{L}})$$

where:

- T_k : Chebyshev polynomials, i.e., $T_0(x) = 1, T_1(x) = x,$
 $T_k(x) = 2xT_{k-1}(x) - T_{k-2}(x)$
- $\tilde{\mathbf{L}}$: Scaled Laplacian, i.e., $\tilde{\mathbf{L}} = 2\lambda_{\max}^{-1}\mathbf{L} - \mathbf{I}$
- **Key property:**
 - A node can influence other nodes within K hops away.
 - Faster computation
- <https://arxiv.org/abs/1606.09375>

From Spectral to Spatial GNNs

Evolution of Graph Convolutional Networks:

1. Spectral GNNs (full eigen)  <https://arxiv.org/abs/1312.6203>
2. ChebNet (polynomial approx)  <https://arxiv.org/abs/1606.09375>
3.  **GCN (*first-order approximation*)**  <https://arxiv.org/abs/1609.02907>
4. Modern spatial GNNs

Graph Convolutional Networks

A Simple Yet Powerful Architecture

Kipf & Welling (2017)

Motivation 🤔

Problems with Previous Approaches:

- Spectral GNNs: Computationally expensive
- ChebNet: Still complex
- Need for simpler, scalable solution

Goals:

- Simple architecture
- Linear complexity
- Good performance

From ChebNet to GCN



ChebNet's Formulation:

$$h_{\theta'} * x \approx \sum_{k=0}^{K-1} \theta_k T_k(\tilde{\mathbf{L}}) x$$

First-Order Approximation:

$$g_{\theta'} * x \approx \theta'_0 x + \theta'_1 (L - I_N) x = \theta'_0 x - \theta'_1 D^{-\frac{1}{2}} A D^{-\frac{1}{2}} x$$

Further Simplification $\theta = \theta_0 = -\theta_1$:

$$g_{\theta} * x \approx \theta (I_N + D^{-\frac{1}{2}} A D^{-\frac{1}{2}}) x$$

Recipe of GCN



- **Step 1:** Add self-loops to the adjacency matrix.

$$A' = A + I_N$$

- **Step 2:** Compute the normalized adjacency matrix.

$$\tilde{A} = \tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}}$$

where $\tilde{D}_{ii} = \sum_j A'$ is the degree matrix of A' .

- **Step 3:** Perform the convolution.

$$Z = \tilde{A}X^{(\ell)}$$

- **Step 4:** Let different dimensions communicate with each other.

$$Z' = ZW^{(\ell)}$$

- **Step 5:** Apply non-linear activation.

$$X^{(\ell+1)} = \sigma(Z')$$


Renormalization Trick 🔧

- GCN is powerful when it is deep (multiple convolutions).
- But, as the depth increases, the training becomes unstable due to **vanishing/exploding gradients**.

Solution: Add self-connections with renormalization:

$$\tilde{A} = A + I_N$$

$$\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$$



Extensions and Variants



1. GraphSAGE

- Sampling-based approach
- Inductive capability

2. GAT

- Attention mechanisms
- Weighted aggregation

3. GIN

- Theoretically more powerful
- WL test equivalent




GraphSAGE



- Previous GNNs are **transductive**: require the entire graph structure during training and cannot generalize to unseen nodes.
- GraphSAGE is **inductive**: can generalize to unseen nodes.

Key idea: Sample a fixed number of nodes within the neighborhood for each node.


- Allow for localized computation for each node.
- Learns transferable patterns




Another Key idea: Aggregation function

- Treat the self node feature and the neighborhood features differently


$$h_v^{(k+1)} = \text{CONCAT} \left(\underbrace{h_v^{(k)}}_{\text{self node feature}}, \underbrace{\text{AGGREGATE}}_{\text{Sum/Mean/Max/LTCM}} \left\{ h_u^{(k)} \mid u \in \mathcal{N}(v) \right\} \right)$$



1. Sample neighborhood



2. Aggregate feature information
from neighbors



3. Predict graph context and label
using aggregated information

Graph Attention Networks (GAT)

Attention Mechanism:

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k \in \mathcal{N}(i)} \exp(e_{ik})}$$

$$e_{ij} = \text{LeakyReLU}(\vec{a}^\top [\underbrace{W h_i^{(k)}}_{\text{self node feature}} , \underbrace{W h_j^{(k)}}_{\text{neighbor node feature}}])$$

Features:

- Learn importance of neighbors
- Multiple attention heads
- Dynamic edge weights

Graph Isomorphism Network (GIN)



Based on Weisfeiler-Lehman Test:

$$h_v^{(k+1)} = \text{MLP}^{(k)} \left((1 + \epsilon^{(k)}) \cdot h_v^{(k)} + \sum_{u \in \mathcal{N}(v)} h_u^{(k)} \right)$$

Key Features:

- Maximally powerful GNNs
- Theoretical connections to graph isomorphism
- Learnable or fixed ϵ

Summary



Key Takeaways:

- GNNs extend CNNs to irregular structures
- Multiple architectures available:
 - GCN: Simple and effective
 - GraphSAGE: Scalable and inductive
 - GAT: Attention-based
 - GIN: Theoretically powerful

Future Directions:

- Scalability
- Expressiveness
- Applications